



WIKIPEDIA
The Free Encyclopedia

Main page
Contents
Featured content
Current events
Random article
Donate to Wikipedia
Wikipedia store

Interaction
Help
About Wikipedia
Community portal
Recent changes
Contact page

Tools
What links here
Related changes
Upload file
Special pages
Permanent link
Page information
Wikidata item
Cite this page

Print/export
Create a book
Download as PDF
Printable version

Languages
Deutsch
Ελληνικά
Español
فارسی
Français
한국어
Italiano
עברית
Magyar
日本語
Norsk bokmål
Polski
Português
Русский
Српски / srpski
Suomi
Українська
Tiếng Việt
中文

Edit links

[Create account](#) [Log in](#)

Article [Talk](#)

[Read](#) [Edit](#) [View history](#)

Topological sorting

From Wikipedia, the free encyclopedia

"Dependency resolution" redirects here. For other uses, see [Dependency \(disambiguation\)](#).

In the field of [computer science](#), a **topological sort** (sometimes abbreviated **toposort**^[1]) or **topological ordering** of a [directed graph](#) is a linear ordering of its [vertices](#) such that for every directed edge *uv* from vertex *u* to vertex *v*, *u* comes before *v* in the ordering. For instance, the vertices of the graph may represent tasks to be performed, and the edges may represent constraints that one task must be performed before another; in this application, a topological ordering is just a valid sequence for the tasks. A topological ordering is possible if and only if the graph has no [directed cycles](#), that is, if it is a [directed acyclic graph](#) (DAG). Any DAG has at least one topological ordering, and [algorithms](#) are known for constructing a topological ordering of any DAG in [linear time](#).

Contents

[\[hide\]](#)


- Examples
- Algorithms
 - 2.1 Application to shortest path finding
- Complexity
- Uniqueness
- Relation to partial orders
- See also
- References
- External links

Examples

[\[edit\]](#)

The canonical application of topological sorting (topological order) is in [scheduling](#) a sequence of jobs or tasks based on their dependencies; topological sorting algorithms were first studied in the early 1960s in the context of the [PERT](#) technique for scheduling in [project management](#) ([Jarnagin 1960](#)). The jobs are represented by vertices, and there is an edge from *x* to *y* if job *x* must be completed before job *y* can be started (for example, when washing clothes, the washing machine must finish before we put the clothes to dry). Then, a topological sort gives an order in which to perform the jobs.

In computer science, applications of this type arise in [instruction scheduling](#), ordering of formula cell evaluation when recomputing formula values in [spreadsheets](#), [logic synthesis](#), determining the order of compilation tasks to perform in [makefiles](#), data [serialization](#), and resolving symbol dependencies in [linkers](#). It is also used to decide in which order to load tables with foreign keys in databases.



The graph shown to the left has many valid topological sorts, including:

- 7, 5, 3, 11, 8, 2, 9, 10 (visual left-to-right, top-to-bottom)
- 3, 5, 7, 8, 11, 2, 9, 10 (smallest-numbered available vertex first)
- 5, 7, 3, 8, 11, 10, 9, 2 (fewest edges first)
- 7, 5, 11, 3, 10, 8, 9, 2 (largest-numbered available vertex first)
- 7, 5, 11, 2, 3, 8, 9, 10 (attempting top-to-bottom, left-to-right)
- 3, 7, 8, 5, 11, 10, 2, 9 (arbitrary)

Algorithms

[\[edit\]](#)

The usual algorithms for topological sorting have running time linear in the number of nodes plus the number of edges, asymptotically, $O(|V| + |E|)$.

One of these algorithms, first described by [Kahn \(1962\)](#), works by choosing vertices in the same order as the eventual topological sort. First, find a list of "start nodes" which have no incoming edges and insert them into a set *S*; at least one such node must exist in a non-empty acyclic graph. Then:

L ← Empty list that will contain the sorted elements

```

S ← Set of all nodes with no incoming edges
while S is non-empty do
    remove a node n from S
    add n to tail of L
    for each node m with an edge e from n to m do
        remove edge e from the graph
        if m has no other incoming edges then
            insert m into S
if graph has edges then
    return error (graph has at least one cycle)
else
    return L (a topologically sorted order)

```

If the graph is a [DAG](#), a solution will be contained in the list L (the solution is not necessarily unique). Otherwise, the graph must have at least one cycle and therefore a topological sorting is impossible.

Reflecting the non-uniqueness of the resulting sort, the structure S can be simply a set or a queue or a stack. Depending on the order that nodes n are removed from set S, a different solution is created. A variation of Kahn's algorithm that breaks ties [lexicographically](#) forms a key component of the [Coffman–Graham algorithm](#) for parallel scheduling and [layered graph drawing](#).

An alternative algorithm for topological sorting is based on [depth-first search](#). The algorithm loops through each node of the graph, in an arbitrary order, initiating a depth-first search that terminates when it hits any node that has already been visited since the beginning of the topological sort or the node has no outgoing edges (i.e. a leaf node):

```

L ← Empty list that will contain the sorted nodes
while there are unmarked nodes do
    select an unmarked node n
    visit(n)

```

```

function visit(node n)
    if n has a temporary mark then stop (not a DAG)
    if n is not marked (i.e. has not been visited yet) then
        mark n temporarily
        for each node m with an edge from n to m do
            visit(m)
        mark n permanently
        unmark n temporarily
        add n to head of L

```

Each node *n* gets *prepended* to the output list L only after considering all other nodes which depend on *n* (all descendants of *n* in the graph). Specifically, when the algorithm adds node *n*, we are guaranteed that all nodes which depend on *n* are already in the output list L: they were added to L either by the recursive call to visit() which ended before the call to visit *n*, or by a call to visit() which started even before the call to visit *n*. Since each edge and node is visited once, the algorithm runs in linear time. This depth-first-search-based algorithm is the one described by [Cormen et al. \(2001\)](#); it seems to have been first described in print by [Tarjan \(1976\)](#).

Application to shortest path finding [\[edit\]](#)

The topological ordering can also be used to quickly compute [shortest paths](#) through a [weighted](#) directed acyclic graph. Let *V* be the list of vertices in such a graph, in topological order. Then the following algorithm computes the shortest path from some source vertex *s* to all other vertices:^[2]

- Let *d* be an array of the same length as *V*; this will hold the shortest-path distances from *s*. Set $d[s] = 0$, all other $d[u] = \infty$.
- Let *p* be an array of the same length as *V*, with all elements initialized to nil. Each $p[u]$ will hold the predecessor of *u* in the shortest path from *s* to *u*.
- Loop over the vertices *u* as ordered in *T*:
 - For each vertex *v* directly following *u* (i.e., there exists an edge from *u* to *v*):
 - Let *w* be the weight of the edge from *u* to *v*.
 - Relax the edge: if $d[v] > d[u] + w$, set
 - $d[v] \leftarrow d[u] + w$,
 - $p[v] \leftarrow u$.

On a graph of n vertices and m edges, this algorithm takes $\Theta(n + m)$, i.e., [linear](#), time.^[2]

Another algorithm explores all edges connecting *to*, rather than *from*, each vertex in V and has the benefit of being an [online algorithm](#): if new vertices are pushed to the end of V over time, the algorithm can output the shortest paths to all vertices in the latest version of V even before V has completely arrived. The algorithm is as follows (differences **bolded**):

- Let d be an array of the same length as V ; this will hold the shortest-path distances from s . Set $d[s] = 0$, all other $d[u] = \infty$.
- Let p be an array of the same length as V , with all elements initialized to `nil`. Each $p[u]$ will hold the predecessor of u in the shortest path from s to u .
- Loop over **the vertices v** as ordered in T :
 - For each **vertex u directly preceding v** (i.e., there exists an edge from u to v):
 - Let w be the weight of the edge from u to v .
 - Relax the edge: if $d[v] > d[u] + w$, set
 - $d[v] \leftarrow d[u] + w$,
 - $p[v] \leftarrow u$.

Complexity [\[edit\]](#)

The [computational complexity](#) of the problem of computing a topological ordering of a directed acyclic graph is **NC²**; that is, it can be computed in $O(\log^2 n)$ time on a [parallel computer](#) using a polynomial number $O(n^k)$ of processors, for some constant k ([Cook 1985](#)). One method for doing this is to repeatedly square the [adjacency matrix](#) of the given graph, logarithmically many times, using [min-plus matrix multiplication](#) with maximization in place of minimization. The resulting matrix describes the longest path distances in the graph. Sorting the vertices by the lengths of their longest incoming paths produces a topological ordering ([Dekel, Nassimi & Sahni 1981](#)).

Uniqueness [\[edit\]](#)

If a topological sort has the property that all pairs of consecutive vertices in the sorted order are connected by edges, then these edges form a directed [Hamiltonian path](#) in the [DAG](#). If a Hamiltonian path exists, the topological sort order is unique; no other order respects the edges of the path. Conversely, if a topological sort does not form a Hamiltonian path, the DAG will have two or more valid topological orderings, for in this case it is always possible to form a second valid ordering by swapping two consecutive vertices that are not connected by an edge to each other. Therefore, it is possible to test in linear time whether a unique ordering exists, and whether a Hamiltonian path exists, despite the [NP-hardness](#) of the Hamiltonian path problem for more general directed graphs ([Vernet & Markenzon 1997](#)).

Relation to partial orders [\[edit\]](#)

Topological orderings are also closely related to the concept of a [linear extension](#) of a [partial order](#) in mathematics.

A partially ordered set is just a set of objects together with a definition of the " \leq " inequality relation, satisfying the axioms of reflexivity ($x \leq x$), antisymmetry (if $x \leq y$ and $y \leq x$ then $x = y$) and [transitivity](#) (if $x \leq y$ and $y \leq z$, then $x \leq z$). A [total order](#) is a partial order in which, for every two objects x and y in the set, either $x \leq y$ or $y \leq x$. Total orders are familiar in computer science as the comparison operators needed to perform [comparison sorting](#) algorithms. For finite sets, total orders may be identified with linear sequences of objects, where the " \leq " relation is true whenever the first object precedes the second object in the order; a comparison sorting algorithm may be used to convert a total order into a sequence in this way. A linear extension of a partial order is a total order that is compatible with it, in the sense that, if $x \leq y$ in the partial order, then $x \leq y$ in the total order as well.

One can define a partial ordering from any DAG by letting the set of objects be the vertices of the DAG, and defining $x \leq y$ to be true, for any two vertices x and y , whenever there exists a [directed path](#) from x to y ; that is, whenever y is [reachable](#) from x . With these definitions, a topological ordering of the DAG is the same thing as a linear extension of this partial order. Conversely, any partial ordering may be defined as the reachability relation in a DAG. One way of doing this is to define a DAG that has a vertex for every object in the partially ordered set, and an edge xy for every pair of objects for which $x \leq y$. An alternative way of doing this is to use the [transitive reduction](#) of the partial ordering; in general, this produces DAGs with fewer edges, but the reachability relation in these DAGs is still the same partial order. By using these constructions, one can use topological ordering algorithms to find linear extensions of partial orders.

See also [\[edit\]](#)

- [tsort](#), a Unix program for topological sorting
- [dep-trace](#) [↗](#) Orders basic dependencies and unfolds nested ones. (basic: without 2D graphical presumption)
- [Feedback arc set](#), a (possibly empty) set of arcs which, if removed from the graph, make it possible to topologically sort it. Useful for dealing with graphs with cycles.
- [D. E. Knuth](#), [The Art of Computer Programming](#), Volume 1, section 2.2.3, which gives an algorithm for topological sorting of a partial ordering, and a brief history.
- [Dependency graph](#)

References [\[edit\]](#)

1.

[^]

^a

^b

Pang, Chaoyi.

"Information Processing Letters" [↗](#).

Science Direct.

Elsevier.

Retrieved 2015-05-21.

2.

[^]

^a

^b

Cormen, Thomas H.; Leiserson, Charles E.; Rivest, Ronald L., Stein, Clifford (2009) [1990].

Introduction to Algorithms (3rd ed.). MIT Press and McGraw-Hill. pp. 655–657. ISBN 0-262-03384-4.

•

Cook, Stephen A. (1985), "A Taxonomy of Problems with Fast Parallel Algorithms", *Information and Control* **64** (1–3): 2–22, doi:10.1016/S0019-9958(85)80041-3 [↗](#).

•

Cormen, Thomas H.; Leiserson, Charles E.; Rivest, Ronald L.; Stein, Clifford (2001), "Section 22.4: Topological sort", *Introduction to Algorithms* (2nd ed.), MIT Press and McGraw-Hill, pp. 549–552, ISBN 0-262-03293-7.

•

Dekel, Eliezer; Nassimi, David; Sahni, Sartaj (1981), "Parallel matrix and graph algorithms", *SIAM Journal on Computing* **10** (4): 657–675, doi:10.1137/0210049 [↗](#), MR 635424 [↗](#).

•

Jarnagin, M. P. (1960), *Automatic machine methods of testing PERT networks for consistency*, Technical Memorandum No. K-24/60, Dahlgren, Virginia: U. S. Naval Weapons Laboratory.

•

Kahn, Arthur B. (1962), "Topological sorting of large networks", *Communications of the ACM* **5** (11): 558–562, doi:10.1145/368996.369025 [↗](#).

•

Tarjan, Robert E. (1976), "Edge-disjoint spanning trees and depth-first search", *Acta Informatica* **6** (2): 171–185, doi:10.1007/BF00268499 [↗](#).

•

Vernet, Oswaldo; Markenzon, Lilian (1997), "Hamiltonian problems for reducible flowgraphs", *Proc. 17th International Conference of the Chilean Computer Science Society (SCCC '97)*, pp. 264–267, doi:10.1109/SCCC.1997.637099 [↗](#).

External links [\[edit\]](#)

- [NIST Dictionary of Algorithms and Data Structures: topological sort](#) [↗](#)
- [Weisstein, Eric W.](#), "TopologicalSort" [↗](#), *MathWorld*.

v · t · e	Sorting algorithms	[hide]
Theory	Computational complexity theory · Big O notation · Total order · Lists · Inplacement · Stability · Comparison sort · Adaptive sort · Sorting network · Integer sorting	
Exchange sorts	Bubble sort · Cocktail sort · Odd–even sort · Comb sort · Gnome sort · Quicksort · Stooge sort · Bogosort	
Selection sorts	Selection sort · Heapsort · Smoothsort · Cartesian tree sort · Tournament sort · Cycle sort	
Insertion sorts	Insertion sort · Shellsort · Splaysort · Tree sort · Library sort · Patience sorting	
Merge sorts	Merge sort · Cascade merge sort · Oscillating merge sort · Polyphase merge sort · Strand sort	
Distribution sorts	American flag sort · Bead sort · Bucket sort · Burstsor · Counting sort · Pigeonhole sort · Proxmap sort · Radix sort · Flashsort	
Concurrent sorts	Bitonic sorter · Batcher odd–even mergesort · Pairwise sorting network	
Hybrid sorts	Block sort · Timsort · Introsort · Spreadsort · JSort	
Other	Topological sorting · Pancake sorting · Spaghetti sort	

Categories: [Graph algorithms](#) | [Sorting algorithms](#) | [Directed graphs](#)