



WIKIPEDIA
The Free Encyclopedia

[Main page](#)

[Contents](#)

[Featured content](#)

[Current events](#)

[Random article](#)

[Donate to Wikipedia](#)

[Wikipedia store](#)

Interaction

[Help](#)

[About Wikipedia](#)

[Community portal](#)

[Recent changes](#)

[Contact page](#)

Tools

[What links here](#)

[Related changes](#)

[Upload file](#)

[Special pages](#)

[Permanent link](#)

[Page information](#)

[Wikidata item](#)

[Cite this page](#)

Print/export

[Create a book](#)

[Download as PDF](#)

[Printable version](#)

Languages

[Eesti](#)

[Français](#)

[한국어](#)

[עברית](#)

[日本語](#)

[Polski](#)

[Português](#)

[Русский](#)

[Српски / srpski](#)

[தமிழ்](#)

[ไทย](#)

[Edit links](#)

[Create account](#) [Log in](#)

Article

[Talk](#)

[Read](#)

[Edit](#)

[View history](#)

XOR swap algorithm

From Wikipedia, the free encyclopedia
(Redirected from [Xor swap algorithm](#))



This article **needs additional citations for verification**. Please help [improve this article](#) by [adding citations to reliable sources](#). Unsourced material may be challenged and removed. *(February 2012)*

In computer programming, the **XOR swap** is an [algorithm](#) that uses the [XOR bitwise operation](#) to [swap](#) values of distinct [variables](#) having the same [data type](#) without using a temporary variable. "Distinct" means that the variables are stored at different memory addresses; the actual values of the variables do not have to be different.

Contents [\[hide\]](#)

- The algorithm
- Proof of correctness
 - Linear algebra interpretation
- Code example
- Reasons for use in practice
- Reasons for avoidance in practice
 - Aliasing
- Variations
- See also
- Notes
- References

x		y	
1010	⊕	0011	= 1001 → x
1001	⊕	0011	= 1010 → y
1001	⊕	1010	= 0011 → x
0011		1010	

Using the XOR swap algorithm to exchange [nibbles](#) between variables without the use of temporary storage

The algorithm [\[edit\]](#)

Conventional swapping requires the use of a temporary storage variable. Using the XOR swap algorithm, however, no temporary storage is needed. The algorithm is as follows:^{[1][2]}

```
X := X XOR Y
Y := X XOR Y
X := X XOR Y
```

The algorithm typically corresponds to three [machine code](#) instructions. Since XOR is a [commutative operation](#), X XOR Y can be replaced with Y XOR X in any of the lines. When coded in assembly language, this commutativity is often exercised in the second line:

Pseudocode	IBM System/370 assembly	x86 assembly
X := X XOR Y	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
Y := X XOR Y	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
X := X XOR Y	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>

In the above System/370 assembly code sample, R1 and R2 are distinct [registers](#), and each XR operation leaves its result in the register named in the first argument. Using x86 assembly, values X and Y are in registers

eax and ebx (respectively), and `xor` places the result of the operation in the first register.

However, the algorithm fails if x and y use the same storage location, since the value stored in that location will be zeroed out by the first XOR instruction, and then remain zero; it will not be "swapped with itself". Note that this is *not* the same as if x and y have the same values. The trouble only comes when x and y use the same storage location, in which case their values must already be equal. That is, if x and y use the same storage location, then the line:

```
X := X XOR Y
```

sets x to zero (because $x = y$ so $X \text{ XOR } Y$ is zero) *and* sets y to zero (since it uses the same storage location), causing x and y to lose their original values.

Proof of correctness [\[edit\]](#)

The [binary operation](#) XOR over bit strings of length N exhibits the following properties (where \oplus denotes XOR):[\[a\]](#)

- **L1. Commutativity:** $A \oplus B = B \oplus A$
- **L2. Associativity:** $(A \oplus B) \oplus C = A \oplus (B \oplus C)$
- **L3. Identity exists:** there is a bit string, 0, (of length M) such that $A \oplus 0 = A$ for any A
- **L4. Each element is its own inverse:** for each A , $A \oplus A = 0$.

Suppose that we have two distinct registers `R1` and `R2` as in the table below, with initial values A and B respectively. We perform the operations below in sequence, and reduce our results using the properties listed above.

Step	Operation	Register 1	Register 2	Reduction
0	Initial value	A	B	—
1	<code>R1 := R1 XOR R2</code>	$A \oplus B$	B	—
2	<code>R2 := R1 XOR R2</code>	$A \oplus B$	$(A \oplus B) \oplus B = A \oplus (B \oplus B) = A \oplus 0 = A$	L2 L4 L3
3	<code>R1 := R1 XOR R2</code>	$(A \oplus B) \oplus A = A \oplus (A \oplus B) = (A \oplus A) \oplus B = 0 \oplus B = B \oplus 0 = B$	A	L1 L2 L4 L1 L3

Linear algebra interpretation [\[edit\]](#)

As XOR can be interpreted as binary addition and a pair of values can be interpreted as a point in two-dimensional space, the steps in the algorithm can be interpreted as 2×2 matrices with binary values. For simplicity, assume initially that x and y are each single bits, not bit vectors.

For example, the step:

```
X := X XOR Y
```

which also has the implicit:

```
Y := Y
```

corresponds to the matrix $\begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix}$ as

$$\begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} x + y \\ y \end{pmatrix}.$$

The sequence of operations is then expressed as:

$$\begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} 1 & 0 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$$

(working with binary values, so $1 + 1 = 0$), which expresses the [elementary matrix](#) of switching two rows (or columns) in terms of the [transvections](#) (shears) of adding one element to the other.

To generalize to where X and Y are not single bits, but instead bit vectors of length n , these 2×2 matrices are replaced by $2n \times 2n$ [block matrices](#) such as $\begin{pmatrix} I_n & I_n \\ 0 & I_n \end{pmatrix}$.

Note that these matrices are operating on *values*, not on *variables* (with storage locations), hence this interpretation abstracts away from issues of storage location and the problem of both variables sharing the same storage location.

Code example [\[edit\]](#)

A C function that implements the XOR swap algorithm:

```
void xorSwap (int *x, int *y) {
    if (x != y) {
        *x ^= *y;
        *y ^= *x;
        *x ^= *y;
    }
}
```

Note that the code does not swap the integers passed immediately, but first checks if their addresses are distinct. This is because, if the addresses are equal, the algorithm will fold to a triple $*x \hat{=} *x$ resulting in zero.

It can be done more easily, without checking if the addresses are equal, using this method:

```
void xorSwap (int *x, int *y) {
    *x = *x ^ *y ^ (*y = *x);
}
```

The XOR swap algorithm can also be defined with a macro:

```
#define XORSWAP(a, b) ((a) ^= (b), (b) ^= (a), (a) ^= (b))
```

Reasons for use in practice [\[edit\]](#)

In most practical scenarios, the trivial swap algorithm using a temporary register is more efficient. Limited situations in which XOR swapping may be practical include:

- on a processor where the instruction set encoding permits the XOR swap to be encoded in a smaller number of bytes
- in a region with high [register pressure](#), it may allow the [register allocator](#) to avoid spilling a register
- in [microcontrollers](#) where available RAM is very limited.

Because these situations are rare, most optimizing compilers do not generate XOR swap code.

Reasons for avoidance in practice [\[edit\]](#)

Most modern compilers can optimize away the temporary variable in the native swap, in which case the native swap uses the same amount of memory and the same number of registers as the XOR swap and is at least as fast, and often faster. The XOR swap is also much less readable and completely opaque to anyone unfamiliar with the technique.

On modern [CPU architectures](#), the XOR technique can be slower than using a temporary variable to do swapping. One reason is that modern CPUs strive to execute instructions in parallel via [instruction pipelines](#). In the XOR technique, the inputs to each operation depend on the results of the previous operation, so they must be executed in strictly sequential order, negating any benefits of [instruction-level parallelism](#).^[3]

Aliasing [\[edit\]](#)

The XOR swap is also complicated in practice by [aliasing](#). As noted above, if an attempt is made to XOR-swap

the contents of some location with itself, the result is that the location is zeroed out and its value lost. Therefore, XOR swapping must not be used blindly in a high-level language if aliasing is possible.

Similar problems occur with [call by name](#), as in [Jensen's Device](#), where swapping `i` and `A[i]` via a temporary variable yields incorrect results due to the arguments being related: swapping via `temp = i; i = A[i]; A[i] = temp` changes the value for `i` in the second statement, which then results in the incorrect lvalue for `A[i]` in the third statement.

Variations [\[edit\]](#)

The underlying principle of the XOR swap algorithm can be applied to any operation meeting criteria L1 through L4 above. Replacing XOR by addition and subtraction gives a slightly different, but largely equivalent, formulation:

```
void addSwap (unsigned int *x, unsigned int *y)
{
    if (x != y) {
        *x = *x + *y;
        *y = *x - *y;
        *x = *x - *y;
    }
}
```

Unlike the XOR swap, this variation requires that the underlying processor or programming language uses a method such as [modular arithmetic](#) or [bignums](#) to guarantee that the computation of `x + y` cannot cause an error due to [integer overflow](#). Therefore, it is seen even more rarely in practice than the XOR swap.

Note, however, that the implementation of `addSwap` above in the C programming language always works even in case of integer overflow, since, according to the C standard, addition and subtraction of unsigned integers follow the rules of [modular arithmetic](#), i. e. are done in the [cyclic group](#) $\mathbb{Z}/2^s\mathbb{Z}$ where s is the number of bits of `unsigned int`. Indeed, the correctness of the algorithm follows from the fact that the formulas $(x + y) - y = x$ and $(x + y) - ((x + y) - y) = y$ hold in any [abelian group](#). This is actually a generalization of the proof for the XOR swap algorithm: XOR is both the addition and subtraction in the abelian group $(\mathbb{Z}/2\mathbb{Z})^s$.

See also [\[edit\]](#)

- [Symmetric difference](#)
- [XOR linked list](#)
- [Feistel cipher](#) (the XOR swap algorithm is a degenerate form of a Feistel cypher)
- [Amiga CD32](#)

Notes [\[edit\]](#)

- [^] The first three properties, along with the existence of an inverse for each element, are the definition of an [Abelian group](#). The last property is a structural feature of XOR not necessarily shared by other Abelian groups.

References [\[edit\]](#)

- [^] ["The Magic of XOR"](#) . Cs.umd.edu. Retrieved 2014-04-02.
- [^] ["Swapping Values with XOR"](#) . graphics.stanford.edu. Retrieved 2014-05-02.
- [^] Amarasinghe, Saman; Leiserson, Charles (2010). "6.172 Performance Engineering of Software Systems, Lecture 2" . MIT OpenCourseWare. Massachusetts Institute of Technology. Retrieved 27 January 2015.

Categories: [Algorithms](#) | [Binary arithmetic](#)

This page was last modified on 8 July 2015, at 06:01.

Text is available under the [Creative Commons Attribution-ShareAlike License](#); additional terms may apply. By using this site, you agree to the [Terms of Use](#) and [Privacy Policy](#). Wikipedia® is a registered trademark of the [Wikimedia Foundation, Inc.](#), a non-profit organization.

[Privacy policy](#) [About Wikipedia](#) [Disclaimers](#) [Contact Wikipedia](#) [Developers](#) [Mobile view](#)

