



WIKIPEDIA  
The Free Encyclopedia

[Main page](#)  
[Contents](#)  
[Featured content](#)  
[Current events](#)  
[Random article](#)  
[Donate to Wikipedia](#)  
[Wikipedia store](#)

[Interaction](#)  
[Help](#)  
[About Wikipedia](#)  
[Community portal](#)  
[Recent changes](#)  
[Contact page](#)

[Tools](#)  
[What links here](#)  
[Related changes](#)  
[Upload file](#)  
[Special pages](#)  
[Permanent link](#)  
[Page information](#)  
[Wikidata item](#)  
[Cite this page](#)

[Print/export](#)  
[Create a book](#)  
[Download as PDF](#)  
[Printable version](#)

[Languages](#)  
[Čeština](#)  
[Español](#)  
[فارسی](#)  
[日本語](#)  
[中文](#)

 [Edit links](#)

[Create account](#) [Log in](#)

[Article](#) [Talk](#)

[Read](#) [Edit](#) [View history](#)

# AA tree


From Wikipedia, the free encyclopedia




This article **needs additional citations for verification**. Please help [improve this article](#) by [adding citations to reliable sources](#). Unsourced material may be challenged and removed. *(June 2011)*

An **AA tree** in [computer science](#) is a form of [balanced tree](#) used for storing and retrieving ordered data efficiently. AA trees are named for [Arne Andersson](#), their inventor.

AA trees are a variation of the [red-black tree](#), a form of [binary search tree](#) which supports efficient addition and deletion of entries. Unlike red-black trees, red nodes on an AA tree can only be added as a right subchild. In other words, no red node can be a left sub-child. This results in the simulation of a [2-3 tree](#) instead of a [2-3-4 tree](#), which greatly simplifies the maintenance operations. The maintenance algorithms for a red-black tree need to consider seven different shapes to properly balance the tree:



An AA tree on the other hand only needs to consider two shapes due to the strict requirement that only right links can be red:



## Contents [\[hide\]](#)

- [1 Balancing rotations](#)
- [2 Insertion](#)
- [3 Deletion](#)
- [4 Performance](#)
- [5 See also](#)
- [6 References](#)
- [7 External links](#)

## Balancing rotations [\[edit\]](#)

Whereas red-black trees require one bit of balancing metadata per node (the color), AA trees require  $O(\log(N))$  bits of metadata per node, in the form of an integer "level". The following invariants hold for AA trees:

1. The level of every leaf node is one.
2. The level of every left child is exactly one less than that of its parent.
3. The level of every right child is equal to or one less than that of its parent.
4. The level of every right grandchild is strictly less than that of its grandparent.
5. Every node of level greater than one has two children.


A link where the child's level is equal to that of its parent is called a *horizontal* link, and is analogous to a red link in the red-black tree. Individual right horizontal links are allowed, but consecutive ones are forbidden; all left horizontal links are forbidden. These are more restrictive constraints than the analogous ones on red-black trees, with the result that re-balancing an AA tree is procedurally much simpler than re-balancing a red-black tree.

Insertions and deletions may transiently cause an AA tree to become unbalanced (that is, to violate the AA tree invariants). Only two distinct operations are needed for restoring balance: "skew" and "split". Skew is a right rotation to replace a subtree containing a left horizontal link with one containing a right horizontal link instead. Split is a left rotation and level increase to replace a subtree containing two or more consecutive right horizontal links with one containing two fewer consecutive right horizontal links. Implementation of balance-preserving insertion and deletion is simplified by relying on the skew and split operations to modify the tree only if needed,

instead of making their callers decide whether to skew or split.


```
function skew is
  input: T, a node representing an AA tree that needs to be rebalanced.
  output: Another node representing the rebalanced AA tree.

  if nil(T) then
    return Nil
  else if nil(left(T)) then
    return T
  else if level(left(T)) == level(T) then
    Swap the pointers of horizontal left links.
    L = left(T)
    left(T) := right(L)
    right(L) := T
    return L
  else
    return T
  end if
end function
```



```
function split is
  input: T, a node representing an AA tree that needs to be rebalanced.
  output: Another node representing the rebalanced AA tree.

  if nil(T) then
    return Nil
  else if nil(right(T)) or nil(right(right(T))) then
    return T
  else if level(T) == level(right(right(T))) then
    We have two horizontal right links. Take the middle node, elevate it, and
    return it.
    R = right(T)
    right(T) := left(R)
    left(R) := T
    level(R) := level(R) + 1
    return R
  else
    return T
  end if
end function
```



## Insertion [\[edit\]](#)

Insertion begins with the normal binary tree search and insertion procedure. Then, as the call stack unwinds (assuming a recursive implementation of the search), it's easy to check the validity of the tree and perform any rotations as necessary. If a horizontal left link arises, a skew will be performed, and if two horizontal right links arise, a split will be performed, possibly incrementing the level of the new root node of the current subtree. Note, in the code as given above, the increment of level(T). This makes it necessary to continue checking the validity of the tree as the modifications bubble up from the leaves.

```

function insert is
  input: X, the value to be inserted, and T, the root of the tree to insert it
  into.
  output: A balanced version T including X.

  Do the normal binary tree insertion procedure. Set the result of the
  recursive call to the correct child in case a new node was created or the
  root of the subtree changes.
  if nil(T) then
    Create a new leaf node with X.
    return node(X, 1, Nil, Nil)
  else if X < value(T) then
    left(T) := insert(X, left(T))
  else if X > value(T) then
    right(T) := insert(X, right(T))
  end if
  Note that the case of X == value(T) is unspecified. As given, an insert
  will have no effect. The implementor may desire different behavior.

  Perform skew and then split. The conditionals that determine whether or
  not a rotation will occur or not are inside of the procedures, as given
  above.
  T := skew(T)
  T := split(T)

  return T
end function

```

## Deletion [\[edit\]](#)

As in most balanced binary trees, the deletion of an internal node can be turned into the deletion of a leaf node by swapping the internal node with either its closest predecessor or successor, depending on which are in the tree or on the implementor's whims. Retrieving a predecessor is simply a matter of following one left link and then all of the remaining right links. Similarly, the successor can be found by going right once and left until a null pointer is found. Because of the AA property of all nodes of level greater than one having two children, the successor or predecessor node will be in level 1, making their removal trivial.

To re-balance a tree, there are a few approaches. The one described by Andersson in his [original paper](#) [↗](#) is the simplest, and it is described here, although actual implementations may opt for a more optimized approach. After a removal, the first step to maintaining tree validity is to lower the level of any nodes whose children are two levels below them, or who are missing children. Then, the entire level must be skewed and split. This approach was favored, because when laid down conceptually, it has three easily understood separate steps:

1. Decrease the level, if appropriate.
2. Skew the level.
3. Split the level.

However, we have to skew and split the entire level this time instead of just a node, complicating our code.

```

function delete is
  input: X, the value to delete, and T, the root of the tree from which it should
  be deleted.
  output: T, balanced, without the value X.

  if nil(T) then
    return T
  else if X > value(T) then
    right(T) := delete(X, right(T))
  else if X < value(T) then
    left(T) := delete(X, left(T))
  else
    If we're a leaf, easy, otherwise reduce to leaf case.
    if leaf(T) then
      return Nil
    else if nil(left(T)) then
      L := successor(T)
      right(T) := delete(value(L), right(T))
      value(T) := value(L)
    else if nil(right(T)) then
      L := predecessor(T)
      left(T) := delete(value(L), left(T))
      value(T) := value(L)
    else
      Internal node, swap with predecessor or successor.
      if nil(left(T)) then
        L := successor(T)
        left(T) := L
        right(T) := delete(value(L), right(T))
      else if nil(right(T)) then
        L := predecessor(T)
        right(T) := L
        left(T) := delete(value(L), left(T))
      else
        Swap with predecessor.
        L := predecessor(T)
        left(T) := L
        right(T) := delete(value(L), right(T))
      end if
    end if
  end if

  return T
end function

```

```

    else
        L := predecessor(T)
        left(T) := delete(value(L), left(T))
        value(T) := value(L)
    end if
end if

Rebalance the tree. Decrease the level of all nodes in this level if necessary, and then skew and split all nodes in the new level.
T := decrease_level(T)
T := skew(T)
right(T) := skew(right(T))
if not nil(right(T))
    right(right(T)) := skew(right(right(T)))
end if
T := split(T)
right(T) := split(right(T))
return T
end function

```

```

function decrease_level is
    input: T, a tree for which we want to remove links that skip levels.
    output: T with its level decreased.

    should_be = min(level(left(T)), level(right(T))) + 1
    if should_be < level(T) then
        level(T) := should_be
        if should_be < level(right(T)) then
            level(right(T)) := should_be
        end if
    end if
    return T
end function

```

A good example of deletion by this algorithm is present in the [Andersson paper](#).

## Performance [\[edit\]](#)

The performance of an AA tree is equivalent to the performance of a red-black tree. While an AA tree makes more rotations than a red-black tree, the simpler algorithms tend to be faster, and all of this balances out to result in similar performance. A red-black tree is more consistent in its performance than an AA tree, but an AA tree tends to be flatter, which results in slightly faster search times.<sup>[1]</sup>


## See also [\[edit\]](#)

- Red-black tree
- B-tree
- AVL tree
- Scapegoat tree

## References [\[edit\]](#)

- ↑ "A Disquisition on The Performance Behavior of Binary Search Tree Data Structures (pages 67-75)"  (PDF).

## External links [\[edit\]](#)

- A. Andersson. [Balanced search trees made simple](#)
- A. Andersson. [A note on searching in a binary search tree](#)
- AA-Tree Applet [by](#) Kubo Kovac
- BSTlib [- Open source AA tree library for C](#) by trijezdci
- AA Visual 2007 1.5 - [OpenSource Delphi program for educating AA tree structures](#)
- Thorough tutorial [Julienne Walker](#) with lots of code, including a practical implementation
- Object Oriented implementation with tests
- A Disquisition on The Performance Behavior of Binary Search Tree Data Structures (pages 67-75)  - Comparison of AA trees, red-black trees, treaps, skip lists, and radix trees

<span>v · t · e</span>	<b>Tree data structures</b>	<span>[show]</span>
<span>v · t · e</span>	<b>Data structures</b>	<span>[show]</span>

Categories: Search trees

This page was last modified on 2 July 2015, at 18:02.

Text is available under the [Creative Commons Attribution-ShareAlike License](#); additional terms may apply. By using this site, you agree to the [Terms of Use](#) and [Privacy Policy](#). Wikipedia® is a registered trademark of the [Wikimedia Foundation, Inc.](#), a non-profit organization.

[Privacy policy](#)
[About Wikipedia](#)
[Disclaimers](#)
[Contact Wikipedia](#)
[Developers](#)
[Mobile view](#)

