


An **ancestry chart** which maps to a perfect depth-4 binary tree.

- A **balanced** binary tree has the minimum possible **maximum height** (a.k.a. depth) for the leaf nodes, because for any given number of leaf nodes the leaf nodes are placed at the greatest height possible.^[clarification needed]



One common balanced tree structure is a binary tree structure in which the left and right subtrees of every node differ in height by no more than 1.^[19] One may also consider binary trees where no leaf is much farther away from the root than any other leaf. (Different balancing schemes allow different definitions of "much farther".^[20])

- A **degenerate** (or **pathological**) tree is where each parent node has only one associated child node.^[citation needed] This means that performance-wise^[clarification needed], the tree will behave like a **linked list** data structure.

Properties of binary trees [edit]

- The number of nodes *n* in a full binary tree, is at least *n* = 2(*h* + 1) − 1 and at most *n* = 2^{*h*+1} − 1, where *h* is the height of the tree. A tree consisting of only a root node has a height of 0.
- The number of leaf nodes *l* in a perfect binary tree, is *l* = (*n* + 1)/2 because the number of non-leaf (a.k.a. internal) nodes *n* − *l* =

∑

k
=
0

log
⁡
2
(
l
)
−
1

2

k

=

2

log
⁡
2
(
l
)

−
1
=
l
−
1

{\displaystyle \sum _{k=0}^{\log _{2}(l)-1}2^{k}=2^{\log _{2}(l)}-1=l-1}

.
- This means that a perfect binary tree with *l* leaves has *n* = 2*l* − 1 nodes.
- In a **balanced** full binary tree, *h* = ⌈log₂(*l*)⌉ + 1 = ⌈log₂((*n* + 1)/2)⌉ + 1 = ⌈log₂(*n* + 1)⌉ (see **ceiling function**).
- In a **perfect** full binary tree, *l* = 2^{*h*} thus *n* = 2^{*h*+1} − 1.
- The maximum possible number of null links (i.e., absent children of the nodes) in a **complete** binary tree of *n* nodes is (*n*+1), where only 1 node exists in bottom-most level to the far left.
- The number of internal nodes in a **complete** binary tree of *n* nodes is ⌊*n*/2⌋.
- For any non-empty binary tree with *n*₀ leaf nodes and *n*₂ nodes of degree 2, *n*₀ = *n*₂ + 1.^[21]

Combinatorics [edit]



This section **does not cite any references or sources**. Please help improve this section by **adding citations to reliable sources**. Unsourced material may be challenged and **removed**. *(July 2014)*

In **combinatorics** one considers the problem of counting the number of full binary trees of a given size. Here the trees have no values attached to their nodes (this would just multiply the number of possible trees by an easily determined factor), and trees are distinguished only by their structure; however the left and right child of any node are distinguished (if they are different trees, then interchanging them will produce a tree distinct from the original one). The size of the tree is taken to be the number *n* of internal nodes (those with two children); the other nodes are leaf nodes and there are *n* + 1 of them. The number of such binary trees of size *n* is equal to the number of ways of fully parenthesizing a string of *n* + 1 symbols (representing leaves) separated by *n* binary operators (representing internal nodes), so as to determine the argument subexpressions of each operator. For instance for *n* = 3 one has to parenthesize a string like *X* * *X* * *X* * *X*, which is possible in five ways:

((*X***X*)**X*)**X*, (*X**(*X***X*))**X*, (*X***X*)*(*X***X*), *X**((*X***X*)**X*), *X**(*X**(*X***X*)).

The correspondence to binary trees should be obvious, and the addition of redundant parentheses (around an already parenthesized expression or around the full expression) is disallowed (or at least not counted as producing a new possibility).

There is a unique binary tree of size 0 (consisting of a single leaf), and any other binary tree is characterized by the pair of its left and right children; if these have sizes *i* and *j* respectively, the full tree has size *i* + *j* + 1. Therefore the number *C_n* of binary trees of size *n* has the following recursive description *C*₀ = 1, and *C_n* =

∑

i
=
0

n
−
1

C

i

C

n
−
1
−
i

{\displaystyle C_{n}=\sum _{i=0}^{n-1}C_{i}C_{n-1-i}}

 for any positive integer *n*. It follows that *C_n* is the **Catalan number** of index *n*.

The above parenthesized strings should not be confused with the set of words of length 2*n* in the **Dyck language**, which consist only of parentheses in such a way that they are properly balanced. The number of such strings satisfies the same recursive description (each Dyck word of length 2*n* is determined by the Dyck subword enclosed by the initial '(' and its matching ')' together with the Dyck subword remaining after that closing parenthesis, whose lengths 2*i* and 2*j* satisfy *i* + *j* + 1 = *n*); this number is therefore also the Catalan number *C_n*. So there are also five Dyck words of length 10:

()()(), ()(()), (())(), (()()), (((()))

These Dyck words do not correspond in an obvious way to binary trees. A bijective correspondence can nevertheless be defined as follows: enclose the Dyck word in an extra pair of parentheses, so that the result can be interpreted as a **Lisp** list expression (with the empty list () as only occurring atom); then the **dotted-pair** expression for that proper list is a fully parenthesized expression (with NIL as symbol and '.' as operator) describing the corresponding binary tree (which is in fact the internal representation of the proper list).

The ability to represent binary trees as strings of symbols and parentheses implies that binary trees can represent the elements of a **free magma** on a singleton set.

Methods for storing binary trees [edit]

Binary trees can be constructed from **programming language** primitives in several ways.

Nodes and references [edit]

In a language with **records** and **references**, binary trees are typically constructed by having a tree node structure which contains some data and references to its left child and its right child. Sometimes it also contains a reference to its unique parent. If a node has fewer than two children, some of the child pointers may be set to a special null value, or to a special **sentinel node**.

This method of storing binary trees wastes a fair bit of memory, as the pointers will be null (or point to the sentinel) more than half the time; a more conservative representation alternative is **threaded binary tree**.^[22]

In languages with **tagged unions** such as **ML**, a tree node is often a tagged union of two types of nodes, one of which is a 3-tuple of data, left child, and right child, and the other of which is a "leaf" node, which contains no data and functions much like the null value in a language with pointers. For example, the following line of code in **OCaml** (an ML dialect) defines a binary tree that stores a character in each node.^[23]

```
type chr_tree = Empty | Node of char * chr_tree * chr_tree
```

Arrays [edit]

Binary trees can also be stored in breadth-first order as an **implicit data structure** in **arrays**, and if the tree is a complete binary tree, this method wastes no space. In this compact arrangement, if a node has an index *i*, its children are found at indices 2*i* + 1 (for the left child) and 2*i* + 2 (for the right), while its parent (if any) is found at index ⌊


i
−
1

2

⌋ (assuming the root has index zero). This

method benefits from more compact storage and better **locality of reference**, particularly during a preorder traversal. However, it is expensive to grow and wastes space proportional to 2^{*h*} · *n* for a tree of depth *h* with *n* nodes.

This method of storage is often used for **binary heaps**. No space is wasted because nodes are added in breadth-first order.



Encodings [edit]

Succinct encodings [edit]

A **succinct data structure** is one which occupies close to minimum possible space, as established by **information theoretical** lower bounds. The number of different binary trees on *n* nodes is **C_n**, the

n th Catalan number (assuming we view trees with identical *structure* as identical). For large n , this is about 4^n ; thus we need at least about $\log_2 4^n = 2n$ bits to encode it. A succinct binary tree therefore would occupy $2n + o(n)$ bits.

One simple representation which meets this bound is to visit the nodes of the tree in preorder, outputting "1" for an internal node and "0" for a leaf. [1] If the tree contains data, we can simply simultaneously store it in a consecutive array in preorder. This function accomplishes this:

```
function EncodeSuccinct(node n, bitstring structure, array data) {
  if n = nil then
    append 0 to structure;
  else
    append 1 to structure;
    append n.data to data;
    EncodeSuccinct(n.left, structure, data);
    EncodeSuccinct(n.right, structure, data);
}
```

The string *structure* has only $2n + 1$ bits in the end, where n is the number of (internal) nodes; we don't even have to store its length. To show that no information is lost, we can convert the output back to the original tree like this:

```
function DecodeSuccinct(bitstring structure, array data) {
  remove first bit of structure and put it in b
  if b = 1 then
    create a new node n
    remove first element of data and put it in n.data
    n.left = DecodeSuccinct(structure, data)
    n.right = DecodeSuccinct(structure, data)
    return n
  else
    return nil
}
```


More sophisticated succinct representations allow not only compact storage of trees but even useful operations on those trees directly while they're still in their succinct form.

Encoding general trees as binary trees [edit]

There is a one-to-one mapping between general ordered trees and binary trees, which in particular is used by Lisp to represent general ordered trees as binary trees. To convert a general ordered tree to binary tree, we only need to represent the general tree in left child-right sibling way. The result of this representation will be automatically binary tree, if viewed from a different perspective. Each node N in the ordered tree corresponds to a node N' in the binary tree; the *left* child of N' is the node corresponding to the first child of N , and the *right* child of N' is the node corresponding to N 's next sibling --- that is, the next node in order among the children of the parent of N . This binary tree representation of a general order tree is sometimes also referred to as a **left child-right sibling binary tree** (LCRS tree), or a **doubly chained tree**, or a **Filial-Heir chain**.

One way of thinking about this is that each node's children are in a **linked list**, chained together with their *right* fields, and the node only has a pointer to the beginning or head of this list, through its *left* field.

For example, in the tree on the left, A has the 6 children {B,C,D,E,F,G}. It can be converted into the binary tree on the right.



The binary tree can be thought of as the original tree tilted sideways, with the black left edges representing *first child* and the blue right edges representing *next sibling*. The leaves of the tree on the left would be written in Lisp as:

((N O) I J) C D ((P) (Q)) F (M))

which would be implemented in memory as the binary tree on the right, without any letters on those nodes that have a left child.

Common operations [edit]

There are a variety of different operations that can be performed on binary trees. Some are **mutator** operations, while others simply return useful information about the tree.

Insertion [edit]

Nodes can be inserted into binary trees in between two other nodes or added after a **leaf node**. In binary trees, a node that is inserted is specified as to which child it is.

Leaf nodes [edit]

To add a new node after leaf node A, A assigns the new node as one of its children and the new node assigns node A as its parent.

Internal nodes [edit]

Insertion on **internal nodes** is slightly more complex than on leaf nodes. Say that the internal node is node A and that node B is the child of A. (If the insertion is to insert a right child, then B is the right child of A, and similarly with a left child insertion.) A assigns its child to the new node and the new node assigns its parent to A. Then the new node assigns its child to B and B assigns its parent as the new node.

Deletion [edit]

Deletion is the process whereby a node is removed from the tree. Only certain nodes in a binary tree can be removed unambiguously.[24]

Node with zero or one children [edit]

Suppose that the node to delete is node A. If A has no children, deletion is accomplished by setting the child of A's parent to **null**. If A has one child, set the parent of A's child to A's parent and set the child of A's parent to A's child.

Node with two children [edit]

In a binary tree, a node with two children cannot be deleted unambiguously.[24] However, in certain binary trees (including **binary search trees**) these nodes can be deleted, though with a rearrangement of the tree structure.


Traversal [edit]

Main article: [Tree traversal](#)


Pre-order, in-order, and post-order traversal visit each node in a tree by recursively visiting each node in the left and right subtrees of the root.

Depth-first order [edit]


In depth-first order, we always attempt to visit the node farthest from the root node that we can, but with the caveat that it must be a child of a node we have already visited. Unlike a depth-first search on graphs, there is no need to remember all the nodes we have visited, because a tree cannot contain cycles. Pre-order is a special case of this. See [depth-first search](#) for more information.



Tree rotations are very common internal operations on self-balancing binary trees.



The process of inserting a node into a binary tree



The process of deleting an internal node in a binary tree

Bibliography [\[edit\]](#)

- Donald Knuth. *The Art of Computer Programming vol 1. Fundamental Algorithms*, Third Edition. Addison-Wesley, 1997. ISBN 0-201-89683-4. Section 2.3, especially subsections 2.3.1–2.3.2 (pp. 318–348).

External links [\[edit\]](#)

- binary trees [↗](#) entry in the FindStat [↗](#) database
- Gamedev.net introduction on binary trees [↗](#)
- Binary Tree Proof by Induction [↗](#)
- Balanced binary search tree on array How to create bottom-up an Ahnentafel list, or a balanced binary search tree on array [↗](#)



v t e	Tree data structures [hide]	
<div>Search trees (dynamic sets/associative arrays)</div>	2–3 · 2–3–4 · AA · (a,b) · AVL · B · B+ · B* · B ^x · (Optimal) Binary search · Dancing · HTree · Interval · Order statistic · (Left-leaning) Red-black · Scapegoat · Splay · T · Treap · UB · Weight-balanced	
<div>Heaps</div>	Binary · Binomial · Fibonacci · Leftist · Pairing · Skew · Van Emde Boas	
<div>Tries</div>	Hash · Radix · Suffix · Ternary search · X-fast · Y-fast	
<div>Spatial data partitioning trees</div>	BK · BSP · Cartesian · Hilbert R · <i>k</i> -d (implicit <i>k</i> -d) · M · Metric · MMP · Octree · Priority R · Quad · R · R+ · R* · Segment · VP · X	
<div>Other trees</div>	Cover · Exponential · Fenwick · Finger · Fusion · Hash calendar · iDistance · K-ary · Left-child right-sibling · Link/cut · Log-structured merge · Merkle · PQ · Range · SPQR · Top	

Categories: Binary trees | Data structures